37 paper claims re-tested on a rig of about 1,500 euros, by verdict

the failures fall into three classes

A the clock

adaptive chunking: one decision
4,751 ms on the edge GPU, budget 298 ms

B a judge against episode length

reward model AUROC 0.818, length alone 0.938
(ArmnetBench; one case)

C learned latents

world model ranks the executed chunk
first in 0.60 of anchors, bar 0.8