37 paper claims re-tested on the 1,350 euro rig, by verdict

the failures fall into three classes

A the clock

adaptive chunking: one decision
4,751 ms on the edge GPU, budget 298 ms

B zero-shot judges

a learned reward model: AUROC 0.818
episode length alone: 0.938

C learned latents

world model ranks the executed chunk
first in 0.60 of anchors, bar 0.8